

STA378 Final Report

Felix Gao

September 2025

1 Abstract

Hyperparameter tuning plays a critical role in determining the efficiency and robustness of optimization solvers used for large-scale, unconstrained problems. We investigate the impact of the L-BFGS memory parameter *mem* on solver performance using a benchmark suite implemented in Julia through the *ISO-Solvers.jl* and *OptimizationProblems.jl* packages. The performance of the algorithm can be measured in term of elapsed time, number of evaluation of the objective function f , its gradient, its Hessian, or the memory used. The existing algorithms highly depend on algorithmic parameters, for instance for L-BFGS:

- *mem*: memory parameter
- r_1 : slope factor in the Wolfe condition when performing the line search
- bk_{max} : maximum number of backtracks when performing the line search

These parameters are currently "guessed" in the literature or the fruit of trial-error experiments even though they are crucial for the performance.

2 Background and Introduction

Unconstrained optimization algorithms

In the field of continuous numerical optimization, the aim is to develop solvers that can locate a local minimum of unconstrained optimization problems:

$$\min_{x \in \mathbb{R}^n} f(x)$$

where $f : \mathbb{R}^n \rightarrow \mathbb{R}$ and using first-order (gradient vector of size n) and eventually using second-order derivatives (Hessian matrix - symmetric matrix of size n) of the objective function f . There exists a variety of algorithms for such problem and you can find a list of implementations in the Julia package JSOSolver.jl. These algorithms are iterative algorithms that start from an initial guess $x_0 \in \mathbb{R}^n$ and compute a follow-up iterate until a stationary point is reached, i.e, $\nabla f(x) \approx 0$.

In Calculus, the most common optimization algorithms are **Gradient Descent** and **Newton's Method**. However, when the problem becomes very large, it can be very inefficient and hardware expensive to locate a local minimum. For example,

3 Methodology and Implementation

4 Results

Dataframe Description

Column	Type	Description
status	Symbol	first_order : solver successfully reached a first-order stationary point; max_time : solver couldn't solve the problem within the time limit; unbounded : the minimum of the objective function goes to $-\infty$.
name	String	Identifier for the optimization problem.
solver	String	Name of the solver/algorithm applied.
mem	Int	Hyperparameter of the solver.
nvar	Int	Number of decision variables in the problem.
time	Float64	Total time(s) for solving the problem.
memory	Float64	Total memory(MB) used for solving the problem.
num_iter	Int	Total number of iterations performed.
nvmops	Int	Number of vector-matrix products.
neval_obj	Int	Number of objective function evaluations.
init_eval_obj_time	Float64	Time (seconds) for the initial objective evaluation.
init_eval_obj_mem	Float64	Memory (MB) for the initial objective evaluation.
init_eval_obj_alloc	Float64	Number of heap allocations for the initial objective evaluation.
neval_grad	Int	Number of gradient evaluations.
init_eval_grad_time	Float64	Time (seconds) for the initial gradient evaluation.
init_eval_grad_mem	Float64	Memory (MB) for the initial gradient evaluation.
init_eval_grad_alloc	Float64	Number of heap allocations for the initial gradient evaluation.
is_init_run	Bool	Indicates whether this is the initial run.

Table 1: Description of columns in the solver benchmark data table.

5 Conclusions

6 Future work